



ML First Project Documentation

Professor: Dr. Kiani

Ilya Shabanpour Fouladi

4003613041

GitHub Rep. link:

https://github.com/ilya-shabanpour/ML_Credit_Prediction

1402 – 1403 second semester

Table of Contents

Foreword	3
Chapter 1: Data Preprocessing	4
1) Useless Columns	4
2) Filling NaN and Unknown Cells	5
3) Dropping The Duplicate Rows	6
4) One Hot Encoding	6
5) Outliers.....	7
1) Outlier Detection	7
2) Outlier Handling	7
6) Train Test Split	8
7) Scaling.....	8
Chapter 2: Model Selection	9
1) Used Models	9
2) Models Comparison	10

Foreword

In this Documentation I want to elaborate on the implementation of my first Machine Learning Project.

Our task is to use a Regression Model to predict the credit limit of clients based on some features that were given to us in the form of a csv dataset.

This goal is achieved through a process that will be explained in two chapters, Data Preprocessing and Model Selection.

In my implementation I've used these Python Libraries:

- Numpy
- Pandas
- Sklearn
- Matplotlib

Chapter 1: Data Preprocessing

In this chapter I'll discuss the process of preprocessing our data so that it will be ready to be given to our models.

At first let's look at our dataset:

	CLIENTNUM	Customer_Age	Gender	Dependent_count	Education_Level	Marital_Status	Income_Category	Card_Category	Months_on_book	Total_Re
0	768805383	45.0	M		3 High School	Married	\$60K - \$80K	Blue	39.0	
1	818770008	49.0	F		5 Graduate	NaN	Less than \$40K	Blue	44.0	
2	713982108	51.0	M		3 Graduate	Married	\$80K - \$120K	Blue	36.0	
3	769911858	40.0	F		4 High School	NaN	Less than \$40K	Blue	34.0	
4	709106358	40.0	M		3 Uneducated	Married	\$60K - \$80K	NaN	21.0	
...
10162	718673358	35.0	M		3 Doctorate	NaN	\$80K - \$120K	Blue	30.0	
10163	715207458	46.0	F		1 Unknown	Single	Less than \$40K	Blue	39.0	
10164	803665983	52.0	M		0 Unknown	NaN	\$60K - \$80K	Blue	46.0	
10165	713183508	39.0	F		1 High School	NaN	Unknown	NaN	36.0	
10166	718893333	52.0	M		3 Graduate	NaN	\$60K - \$80K	Blue	36.0	

Credit_Limit	Total_Revolving_Bal	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	Total_Trans_Ct	Total_Ct_Chng_Q4_Q1	Avg_Utilization_Ratio	Unnamed: 19
12691.0	777	1.335	1144	42	1.625	0.061	NaN
8256.0	864	1.541	1291	33	3.714	0.105	NaN
3418.0	0	2.594	1887	20	2.333	0.000	NaN
3313.0	2517	1.405	1171	20	2.333	0.760	NaN
4716.0	0	2.175	816	28	2.500	0.000	NaN
...
13590.0	1528	0.728	2137	52	0.486	0.112	NaN
2029.0	1074	0.514	4802	90	0.800	0.529	NaN
2742.0	2184	0.592	3829	72	0.532	0.796	NaN
2751.0	1158	0.821	4861	82	0.822	0.421	NaN
18336.0	2077	0.718	7509	84	0.647	0.113	NaN

As you can see there are useless columns, NaN cells, Unknown cells, and both numerical and String type data.

1) Useless Columns

The first column of our dataset is the client number which has no meaning for our prediction of client's credit limit. The last column also contains NaN cell for each row, so the aforementioned columns should be dropped and not be considered as features:

```
# drop last column because all NaN
df.drop("Unnamed: 19", axis=1, inplace=True)
df.drop("CLIENTNUM", axis=1, inplace=True)
```

2) Filling NaN and Unknown Cells

NaN and Unknown Cells must be Handled as they are one of the important parts of our data preprocessing. One way to handle this event is to drop the rows that contain such elements, but this is undesired as it causes the removal of nearly half of the dataset. So, we should take another approach.

One way is to calculate the mode or mean of that column and replace any Unknown or NaN values with mean or mode of that column.

For string type columns mode is used and for numerical type columns mean of that column.

Here is the code:

```
# replace mod of Marital Status for NaN and Unknown
mod_marriage = df["Marital_Status"].mode()[0]
new_col = df["Marital_Status"].replace(np.nan, mod_marriage)
df["Marital_Status"] = new_col.values
new_col = df["Marital_Status"].replace("Unknown", mod_marriage)
df["Marital_Status"] = new_col.values

# replace mod of Gender for NaN
mod_gender = df["Gender"].mode()[0]
new_col = df["Gender"].replace(np.nan, mod_gender)
df["Gender"] = new_col.values

# replace mod of Education_Level for Unknown
mod_edu = df["Education_Level"].mode()[0]
new_col = df["Education_Level"].replace("Unknown", mod_edu)
df["Education_Level"] = new_col.values

# replace mod Income_Category for Unknown
mod_inc = df["Income_Category"].mode()[0]
new_col = df["Income_Category"].replace("Unknown", mod_inc)
df["Income_Category"] = new_col.values

# replace mod Card_Category for NaN
mod_card_cat = df["Card_Category"].mode()[0]
new_col = df["Card_Category"].replace(np.nan, mod_card_cat)
df["Card_Category"] = new_col.values

# replace mean Months_on_book for NaN
mean_month = df["Months_on_book"].mean().round()
new_col = df["Months_on_book"].replace(np.nan, mean_month)
df["Months_on_book"] = new_col.values

# replace mean Total_Relationship_Count for NaN
mean_total = df["Total_Relationship_Count"].mean().round()
new_col = df["Total_Relationship_Count"].replace(np.nan, mean_total)
df["Total_Relationship_Count"] = new_col.values
```

3) Dropping The Duplicate Rows

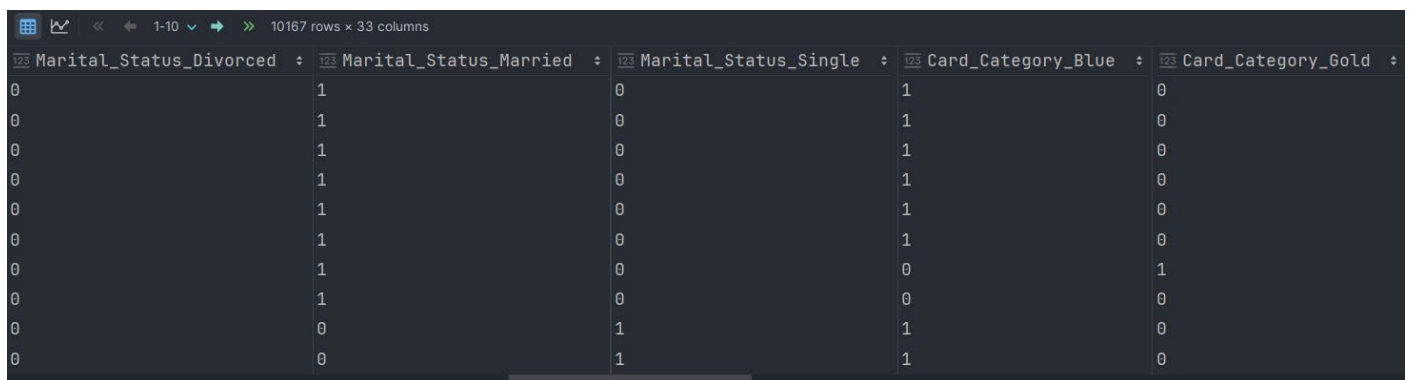
Now that NaN and Unknown cells are filled, some rows might become duplicated. In order to fix that we use the following code to drop those rows:

```
df = df.drop_duplicates()
```

4) One Hot Encoding

It's time to convert string type elements to numbers. In order to do that we can use a method called "One Hot Encoding".

In this method each unique element is considered to be a separate feature and column in our dataset. If an element is in a row, then the value in that column is 1 and 0 if not. Everything will be clearer with looking at an output of this method:



The screenshot shows a Jupyter Notebook interface with a DataFrame containing 10167 rows and 33 columns. The visible columns are Marital_Status_Divorced, Marital_Status_Married, Marital_Status_Single, Card_Category_Blue, and Card_Category_Gold. The data is as follows:

	Marital_Status_Divorced	Marital_Status_Married	Marital_Status_Single	Card_Category_Blue	Card_Category_Gold
0	0	1	0	1	0
0	0	1	0	1	0
0	0	1	0	1	0
0	0	1	0	1	0
0	0	1	0	1	0
0	0	1	0	1	0
0	0	1	0	0	1
0	0	1	0	0	0
0	0	0	1	1	0
0	0	0	1	1	0

Here is the code for implementing this method:

```
# Now we convert the string type columns into numbers using One Hot Encoding
df_encoded = pd.get_dummies(df['Marital_Status'], prefix='Marital_Status')
df_encoded = df_encoded.astype(int)
df = pd.concat([df, df_encoded], axis=1)
df.drop('Marital_Status', axis=1, inplace=True)

df_encoded = pd.get_dummies(df['Card_Category'], prefix='Card_Category')
df_encoded = df_encoded.astype(int)
df = pd.concat([df, df_encoded], axis=1)
df.drop('Card_Category', axis=1, inplace=True)

df_encoded = pd.get_dummies(df['Gender'], prefix='Gender')
df_encoded = df_encoded.astype(int)
df = pd.concat([df, df_encoded], axis=1)
df.drop('Gender', axis=1, inplace=True)

df_encoded = pd.get_dummies(df['Education_Level'], prefix='Education_Level')
df_encoded = df_encoded.astype(int)
```

```
df = pd.concat([df, df_encoded], axis=1)
df.drop("Education_Level", axis=1, inplace=True)

df_encoded = pd.get_dummies(df["Income_Category"], prefix="Income_Category")
df_encoded = df_encoded.astype(int)
df = pd.concat([df, df_encoded], axis=1)
df.drop("Income_Category", axis=1, inplace=True)
```

5) Outliers

This is a very important step in our process. It consists of two parts:

- Outlier Detection
- Outlier Handling

1) Outlier Detection

For detecting outliers, I have come across two methods. The first one requires assuming that our data has Normal distribution. Then we can use the interval of $[\mu - 3 * \sigma, \mu + 3 * \sigma]$ to detect any cell in a column that is outside of this range (the mean and std. are calculated for each column).

The other method is for skewed distributions. We should first find the first and third Quartiles (Q1 and Q3) of each column and find the IQR (Interquartile range) that is:

$$\text{IQR} = Q3 - Q1$$

And then find the interval of $[Q1 - 1.5 * \text{IQR}, Q3 + 1.5 * \text{IQR}]$. Again, any cell in a column outside of this range is considered an outlier.

2) Outlier Handling

After detecting outliers it's time to handle them. There are two approaches:

1 – dropping rows containing outliers:

This method is not desired because we may lose a large part of our data.

2 – replacing outliers with mean value of that column

I have used IQR and mean value replacement in my process.

Here is the code:

```

result = pd.DataFrame(index=data.index, columns=data.columns)

# Loop through each column in the DataFrame
for col in data.columns:
    Q1 = data[col].quantile(0.25) # 25th percentile (Q1)
    Q3 = data[col].quantile(0.75) # 75th percentile (Q3)
    IQR = Q3 - Q1 # Interquartile Range (IQR)
    lower_bound = Q1 - 1.5 * IQR # Lower bound for outliers
    upper_bound = Q3 + 1.5 * IQR # Upper bound for outliers

    # Replace outliers with mode for the current column
    mean_val = data.loc[(data[col] >= lower_bound) & (data[col] <=
upper_bound), col].mean()
    result[col] = data[col].apply(lambda x: mean_val if x < lower_bound or x
> upper_bound else x)

```

6) Train Test Split

Now I split the data into test and train.

Train is used to train the model to find the best Regressor.

Test is for checking the MSE of our Regressor.

I have set 10% of the data for test and the rest for train.

Shuffling the data ensures that you have difficult data in your train batch.

```

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.1,
shuffle=True)

```

7) Scaling

Scaling our data makes it reasonable for our model. Columns have values in different ranges, scaling the data ensures that our data is in a specified range.

I've tested two scaling methods:

1 - Min Max Scaling

2- Standard Scaling

The outcome of both Scaling methods didn't differ much in the final test and I decided to use Min Max Scaling in the end:

```

scaler = MinMaxScaler()

x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)

```


Chapter 2: Model Selection

1) Used Models

I have used these models:

1 – Linear Regression

2 – Polynomial Ridge Regression

3 – MLP (Multiple Layer Perceptron)

4 – Random Forest

```
lr = LinearRegression()
ridge = Ridge()
rf = RandomForestRegressor(n_estimators=100, max_depth=15, random_state=10)
mlp = MLPRegressor(max_iter=5000, random_state=0, learning_rate_init=0.005)

rf.fit(x_train, y_train)
y_pred_rf = rf.predict(x_test)

lr.fit(x_train, y_train)
y_pred_lr = lr.predict(x_test)

ridge.fit(x_train_poly, y_train)
y_pred_ridge = ridge.predict(x_test_poly)

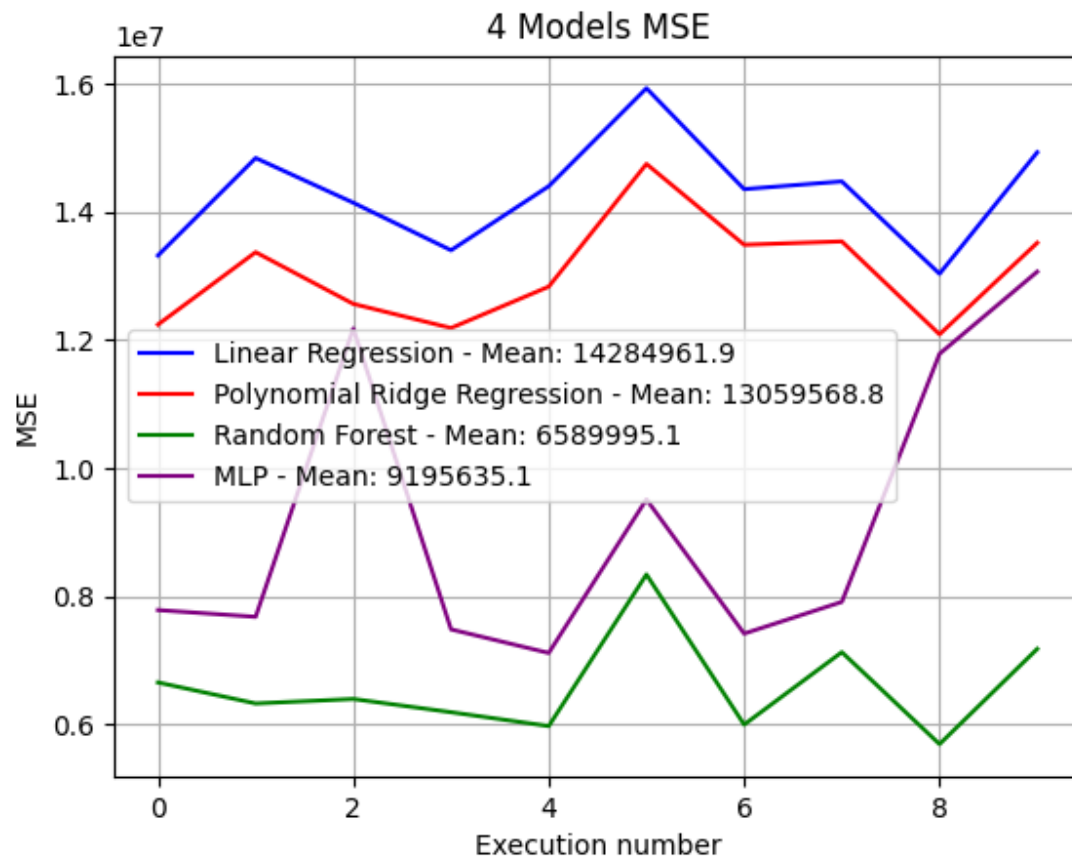
mlp.fit(x_train, y_train)
y_pred_mlp = mlp.predict(x_test)

mse_rf = round(mean_squared_error(y_true=y_test, y_pred=y_pred_rf))
mse_lr = round(mean_squared_error(y_true=y_test, y_pred=y_pred_lr))
mse_ridge = round(mean_squared_error(y_true=y_test, y_pred=y_pred_ridge))
mse_mlp = round(mean_squared_error(y_true=y_test, y_pred=y_pred_mlp))
```

2) Models Comparison

After fitting with every model and calculating the MSE for 10 runs, I calculated the mean MSE of every model and plotted it using Matplotlib library.

Here are my results:



As you can see Random Forest has the lowest mean MSE compared to the other models. The mean MSE of Random Forest is the range of 6 million and it is as low as I could get MSE to be in any of used models.

So, my final model is Random Forest and the hyperparameters were set to be:

`n_estimators = 100` (determines the number of different decision trees used for averaging, and also determines the different subsamples used)

`max_depth = 15` (the maximum depth of the decision trees in the forest)